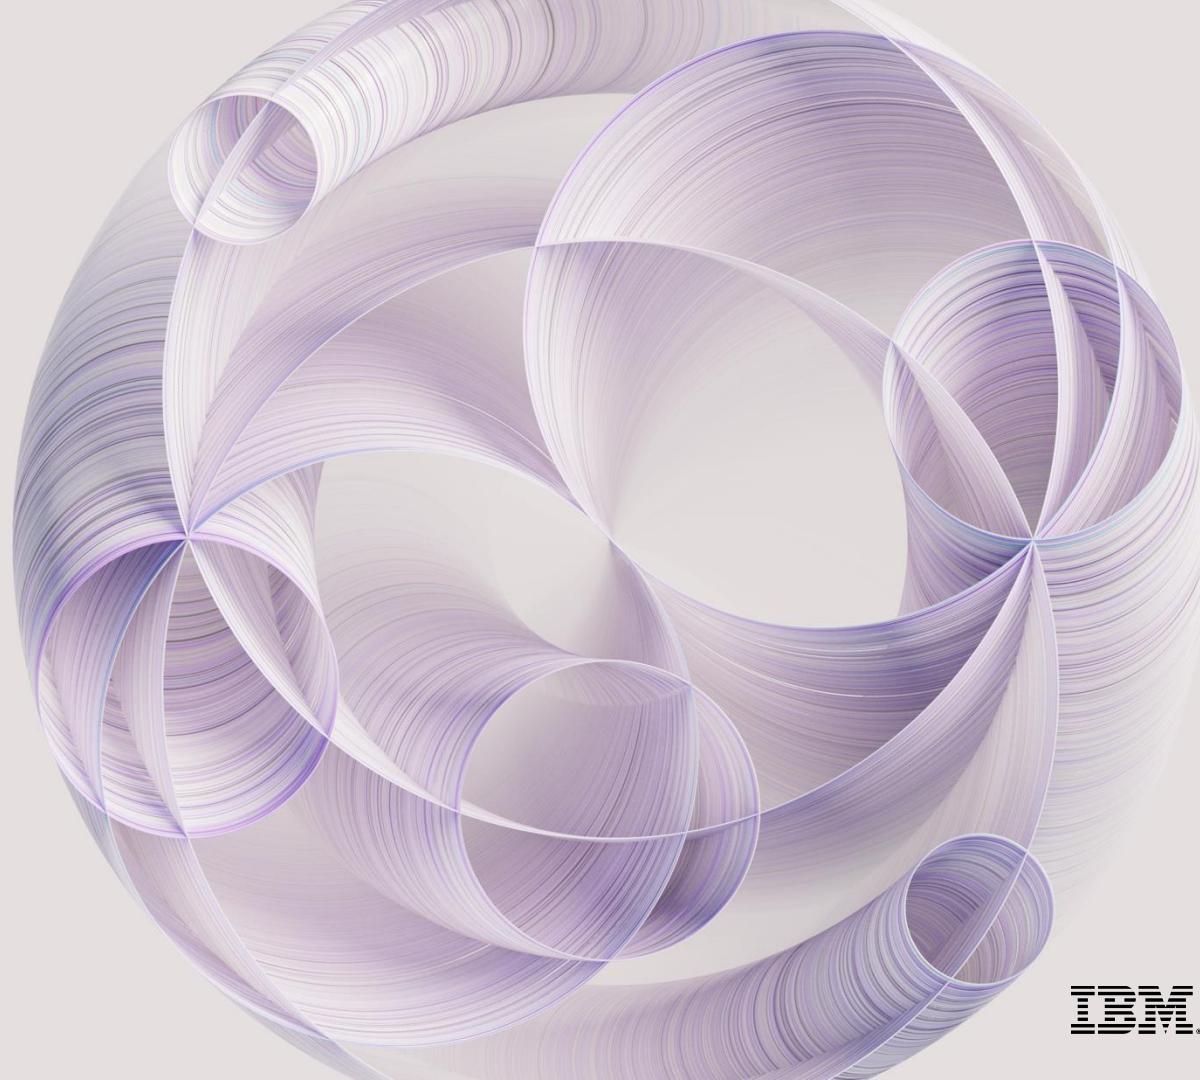


Streamline the
development of
AI applications
with [watsonx.ai](#)



Hybrid Cloud

Unify on-prem, public, private clouds and edge to scale virtualization and AI across environments



Transaction Processing

Deliver unmatched transactional performance, security and reliability



Automation

Automate technology lifecycle management with AI for productivity, resiliency and spend optimization



Data

Access trusted and secure data to drive AI productivity



AI

Open and Trusted

Data

Access trusted and secure data to drive
AI productivity



AI Productivity

Reinvent how work is done with AI agents/assistants

– AI Assistants

AI/ML Ops

Work with AI models, tools and governance that's built for business—engineered to ensure trust and scalability in applications

– AI Models
– AI Tools
– AI Governance

Data Fabric

Bring all your business data together and optimize how it moves through your systems to scale analytics and AI in your applications while protecting it.

– Databases
– Data Intelligence
– Data Integration
– Data Security

Data Storage

Store data across edge, core and clouds

– Software-defined Storage

Data

Access trusted and secure data to drive
AI productivity



AI Productivity

Reinvent how work is done with AI agents/assistants

AI/ML Ops

Work with AI models, tools and governance that's built for business—engineered to ensure trust and scalability in applications

Data Fabric

Bring all your business data together and optimize how it moves through your systems to scale analytics and AI in your applications while protecting it

Data Storage

Store data across edge, core and clouds

AI Assistants



watsonx Code Assistant™



watsonx Orchestrate™



Planning Analytics

AI Models



Granite™

AI Tools



watsonx.ai™

AI Governance



watsonx.governance™

Databases



watsonx.data™

Data Intelligence



Data Product Hub



Knowledge Catalog



Manta Data Lineage

Data Integration



DataStage®



Databand®



Streamsets

Data Security



Guardium® Data Security Center

Software-defined Storage



Storage Ceph®

watsonX

A portfolio of AI products that accelerates the impact of generative AI in core workflows to drive productivity.

watsonX.ai

Enterprise-grade AI studio that helps AI builders innovate with all the APIs, tools, models, and runtimes to build AI solutions

Featuring **IBM Granite**, and popular third-party models including **Mixtral**, **Llama** series

watsonX.data

The **hybrid, open data lakehouse** to power AI and analytics with all your data, anywhere

watsonX.governance

End-to-end toolkit for AI governance to manage **risk and compliance across the entire AI lifecycle**.

watsonX Orchestrate

An enterprise-ready solution that helps create, deploy, and manage AI assistants and agents to automate processes and workflows.

watsonX Code Assistant

Accelerate development, **application modernization**, and **assist with IT Operations**

Why IBM Watsonx for scaling enterprise AI to drive productivity

Open

- Offers choice to train the right foundation models, including open-source models, and the choice of data, tools, and frameworks to achieve desired business outcomes.
- Run AI wherever the business needs to, across any cloud, at scale.

Trusted

- Built with open and transparent technology to give enterprises confidence in their AI and meet regulatory compliance demands.
- Responsible AI and protected data backed by enterprise governance and security controls.

Integrated

- Integrates technology seamlessly into existing infrastructures, systems, and processes with choice of cloud to transform the enterprise and drive productivity from within.
- Embedded AI for targeted use cases that drives enterprise scale productivity.

Streamline the development of AI applications with watsonx.ai



Model selection



Model customization



AI development and deployment



AI governance

The screenshot shows the IBM WatsonX developer interface. At the top, there's a navigation bar with 'IBM watsonx' on the left, and 'Upgrade', 'IBM', 'Dallas', and a user profile icon on the right. Below the header, a welcome message 'Welcome back, Laura' is displayed. A main banner features the text 'Train, validate, tune and deploy AI models.' and a 'Customize my journey' button. To the right of the banner are three cards: 'Chat and build prompts with foundation models' (with a 'Start chatting...' input field and 'Open Prompt Lab' button), 'Tune a foundation model with labeled data with Tuning Studio', and 'Work with data and models in Python or R notebooks with Jupyter notebook editor'. Below this is a 'Discover' section with a 'Resource hub' table:

Foundation models	Prompts	→
Explore foundation models from IBM and other third-parties depending on your use case.	Data	→
Explore foundation models →	Projects	→
	Notebooks	→

On the right side of the 'Discover' section, there's a 'What's new' sidebar with three items:

- New watsonx Developer Hub to start coding fast Oct 20, 2024 →
- IBM watsonx.governance is available in the Sydney region Dec 08, 2024 →
- Deploy foundation models on demand in the Dallas region Dec 05, 2024 →

The enterprise-grade AI developer studio

Integrated AI platform with all the APIs, tools, models, and runtimes to simplify and scale the development and deployment of AI

Applications, Agents, Assistants

Developer toolkit

Tools

- API, SDK, CLI
- AutoAI RAG
- Tuning Studio
 - InstructLab *private preview*
- Prompt Lab
 - Chat with documents/images
 - Customize model parameters
- Developer Hub
- Build time model evaluation

Deployment

- AI services for agents and models

Frameworks

- Third-party data sources and tooling
- Open-source frameworks (Node.js SDK, LangChain, LlamaIndex, CrewAI)
- Flows Engine

Governance

Common core services

Project asset management, jobs, user access, data connectors

AI runtimes

AI Infrastructure

- Red Hat OpenShift AI
- vLLM

Model choices

- IBM Granite™
- Third party (e.g. Mistral AI, Meta)
- Open source
- Custom foundation models

Machine learning

Predictive and prescriptive analytics

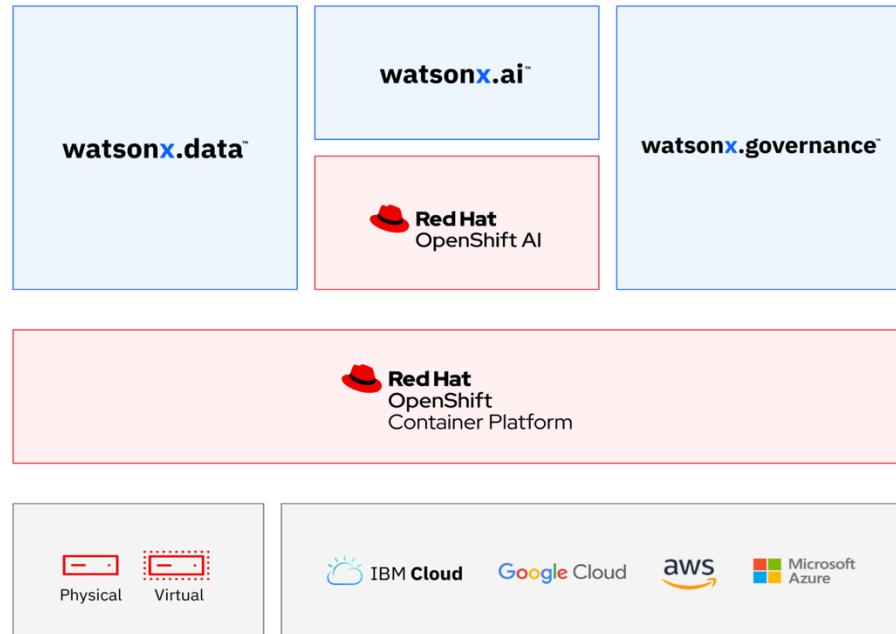
Hybrid cloud platform

watsonx.ai and Red Hat

Watsonx.ai integrated with **Red Hat OpenShift AI** & **Red Hat OpenShift Container Platform** delivers an accelerated time to value through an optimized generative AI and ML workflow.

Key benefits:

- A complete platform for managing AI workloads, training models and scaling AI deployments
- Access to high quality, governed, and trusted AI through curated IBM-developed, third-party, and open-source foundation models
- Deploy to on-premises environments, private or public clouds
- Integrations to the wider watsonx and Cloud Pak for Data platforms.



Creating Generative AI Solutions

Model Selection

watsonx.ai Provided Models

A highly curated library with IBM-built, open-source and third-party foundation models that are pre-deployed and ready to use in watsonx.ai for a wide range of business requirements and budgetary considerations. These models are deployed on multi-tenant GPU's.

Key benefits:

- Great for experimentation and prototyping to determine which model is best for your use case
- Pay only for the tokens that are consumed, no time commitment
- Ideal for clients that need periodic or infrequent use of foundation models
- Indemnification offered for IBM and select models

Learn more in [documentation](#)

Most models available in all regions

The screenshot shows the IBM WatsonX interface with a search bar at the top. Below it, a section titled "Select a foundation model" displays a grid of 12 model cards. Each card includes the model name, provider, type, and a brief description. The models listed are: granite-13b-chat-v2, granite-3-1-2b-instruct, granite-3-8b-instruct, granite-3-2b-instruct, granite-3-1-8b-instruct, llama-3-3-70b-instruct, granite-guardian-3-8b, mistral-large, granite-guardian-3-2b, pixtral-12b, granite-20b-multilingual, and codestral-22b. The "granite-13b-chat-v2" card has a blue circular badge with the number "9".

Model	Provider	Type	Description
granite-13b-chat-v2	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
granite-3-1-2b-instruct	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
granite-3-8b-instruct	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
granite-3-2b-instruct	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
granite-3-1-8b-instruct	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
llama-3-3-70b-instruct	Meta	Provided	This version of Llama-3-3-70b-instruct is also the FP8 quantized version of the original FP16 weights.
granite-guardian-3-8b	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
mistral-large	Mistral AI	Provided	Mistral Large 240T, the most advanced Large Language Model (LLM) developed by Mistral AI, is exceptionally powerful.
granite-guardian-3-2b	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
pixtral-12b	Mistral AI	Provided	The Pixtral-12b is a Multimodal Model of 12B parameters plus a 400M parameter vision encoder.
granite-20b-multilingual	IBM	Provided	The Granite model series is a family of IBM-trained, dense decoder-only models, which are particularly well-suited for generative...
codestral-22b	Mistral AI	Provided	Codestral-22b is an open-weight generative AI model explicitly designed for code generation tasks.

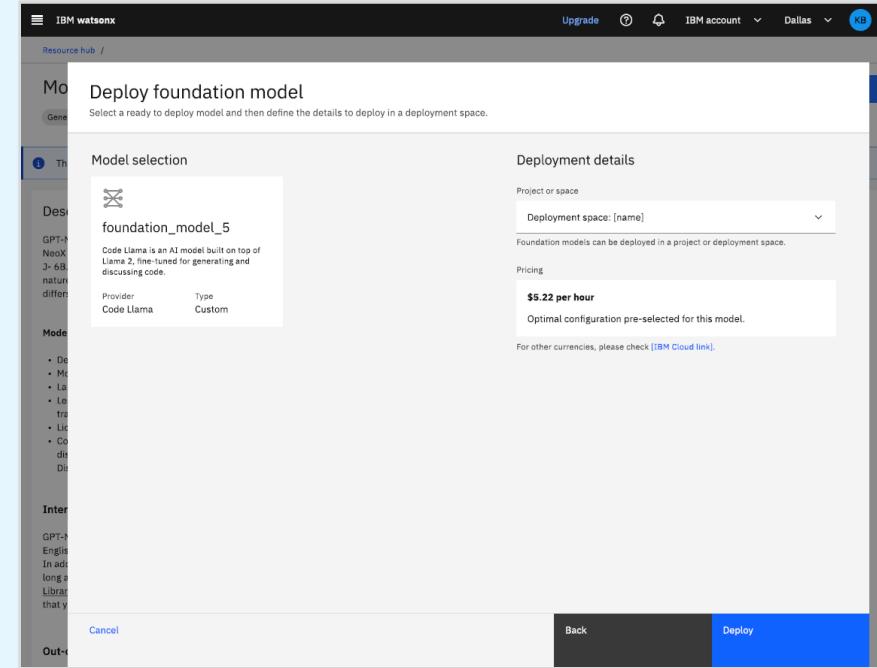
Deploy on Demand Models

A curated collection of high quality and popular foundation models that can be easily deployed to a dedicated GPU with just a few clicks.

Key benefits:

- Faster interactions with a dedicated, single-tenant deployment, hosted 24/7 until de-provisioned by client
- Predictable hourly hosting price compared to variable token-based pricing
- Supports full context length of the model
- No rate limits on the inference requests per second (vs. 8 req/s limit on pre-deployed models)

Learn more in [documentation](#)



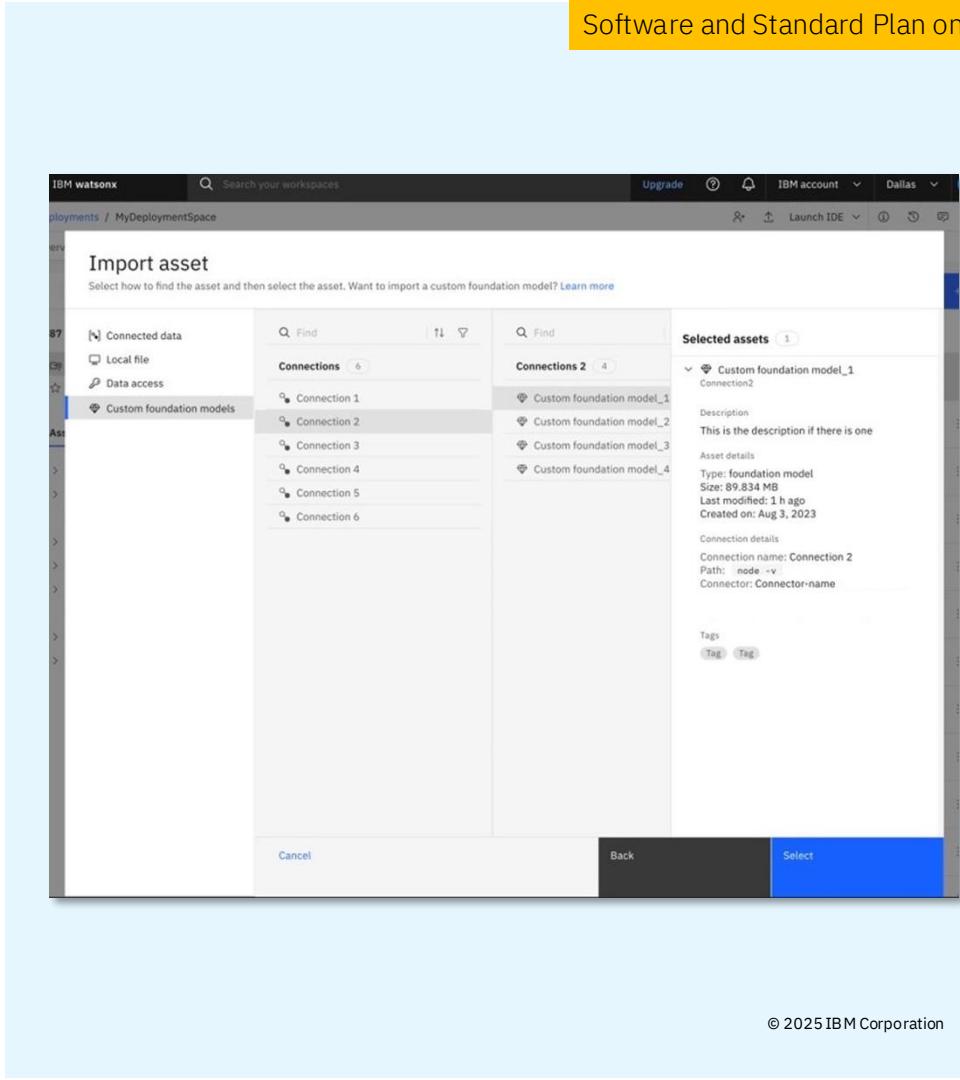
Custom Foundation Models

Import custom foundation models for maximum flexibility in how generative AI solutions are developed.

Key benefits:

- Leverage externally fine-tuned LLMs that support a specific language or are customized for an industry or business domain
- Import a model that is not already provided by IBM.
- Compatible with 1000's of models from repositories like Hugging Face, which provides access to a huge selection of open-source foundation models

Learn more in the [blog](#), [tutorial](#) or [documentation](#)



Access 1000+ models in watsonx.ai from a mix of platform-provided models and custom foundation models available for import

	Provided Models (Ready to use models)	Deploy on Demand	Import Custom Foundation Models	Software
Overview	Multi-tenant SaaS offering where capacity is shared	Single-tenant SaaS offering with curated set of models from IBM	Single-tenant SaaS offering with customer-provided models	On-premises deployment on customer infrastructure
Cost Structure	Token-based pricing, billed on both input and output tokens; available for all SaaS Plans	Hourly pricing - includes hosting and inference; billed while model is deployed; Standard Plan only	Hourly pricing - includes hosting and inference; billed while model is deployed; Standard Plan only	No additional model pricing; client is billed on software entitlement
Benefits	<ul style="list-style-type: none"> Pay only for the tokens that are consumed Indemnity offered for IBM and select models 	<ul style="list-style-type: none"> Predictable cost based on hourly model deployment Dedicated GPU capacity Ease of deployment Indemnity offered for IBM and select models 	<ul style="list-style-type: none"> Predictable cost based on hourly model deployment Dedicated GPU capacity 	<ul style="list-style-type: none"> Maximum flexibility in choosing which models to deploy and on what kind of GPUs Air-gapped environment for secure workloads
Drawbacks	<ul style="list-style-type: none"> GPU capacity is shared; unable to take full advantage of the throughput 	<ul style="list-style-type: none"> Deprovision the model when not in use to optimize costs 	<ul style="list-style-type: none"> Requires user to import relevant model files Deprovision the model when not in use to optimize costs 	<ul style="list-style-type: none"> Infrastructure setup required
Best Fit for	<ul style="list-style-type: none"> Experimentation Small companies with infrequent use 	<ul style="list-style-type: none"> Large production workloads Enterprises with SLA requirements 	<ul style="list-style-type: none"> Large production workloads Enterprises with SLA requirements 	<ul style="list-style-type: none"> Large production workloads Enterprises with SLA and security requirements

Model Customization

watsonx.ai Tuning Studio

Improve foundation model performance through customization using a variety of methods:

- Fine Tuning
- LoRA/QLoRA (PEFT)

Easy to use, step by step interface to accelerate the setup time.

Customized models can be deployed within watsonx.ai and inferenced through the Prompt Lab or API.

Learn more in [documentation](#)

The screenshot shows the IBM Watson Tuning Studio interface. At the top, there's a navigation bar with 'IBM Watsonx', a search bar, and various account and location options like 'Upgrade', 'IBM account', 'Dallas', and a battery icon. Below the header, the page title is 'Projects / Kate's Sandbox / New tuned model'. A timestamp 'Last saved: 00:00:00 PM' is visible. The main content area is titled 'Customize tuned model' and 'New tuned model'. It has two main sections: 'Configure details' and 'Submission details'. In 'Configure details', there's a section for 'Which foundation model do you want to tune?' with a dropdown set to 'granite-13b-instruct-v2'. Below it, there are three options: 'Prompt tuning' (disabled), 'Fine tuning' (disabled), and 'InstructLab' (selected). There's also a section for 'How do you want to work?' with 'No code' selected. In 'Submission details', there's a 'Submit for training' section with a note about reviewers checking grounding data, a 'Github ID' field containing 'jane.doe', and a note about GitHub ID lookup. Below that, there are two cards for 'New skill: Arithmetic_w_grammar': one for teaching grammar and arithmetic, and another identical one below it. At the bottom right, there's a blue 'Submit for review' button.

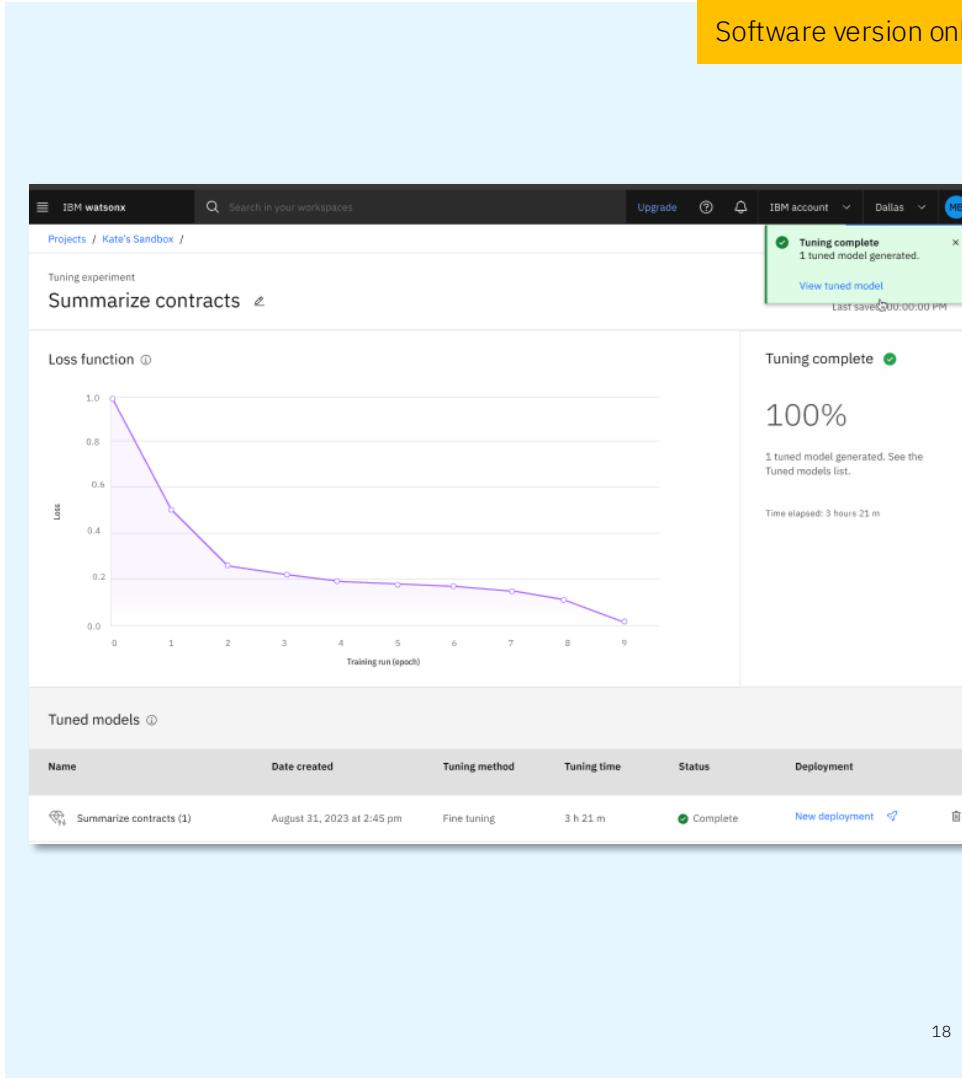
Tuning Studio: Fine Tuning

Unlock the full potential of foundation models by customizing it for a specific task, use case or business domain with structured data.

Key benefits:

- Improved performance and accuracy
- Reduced Bias from original training data
- Easy to use interface that supports structured training data in files or connected data sources
- Adjust parameters and view performance metrics

Learn more in [documentation](#)



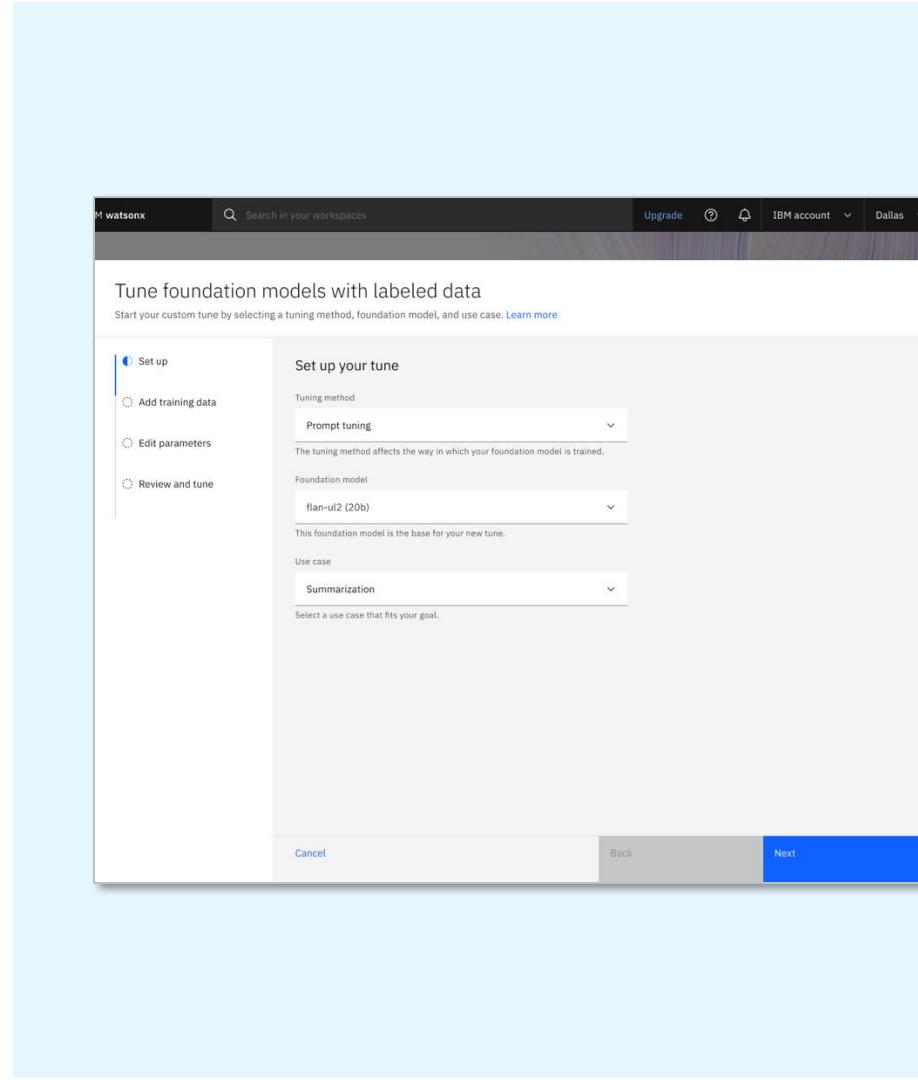
Tuning Studio: LoRA/QLoRA Tuning

Leverage LoRA and QLoRA to further train a model with domain data. This method drastically reduces memory, computation, and time compared to full fine-tuning.

Key benefits:

- Efficient, low-cost way of adapting an AI foundation model to new tasks
- Tune the model by changing only a subset of model parameters (not the entire base model)
- Achieve comparable results to full fine-tuning depending on base model size and data used
- Minimize the risk of catastrophic forgetting

This method is a subset of Parameter Efficient Fine Tuning (PEFT). Learn more in [documentation](#).



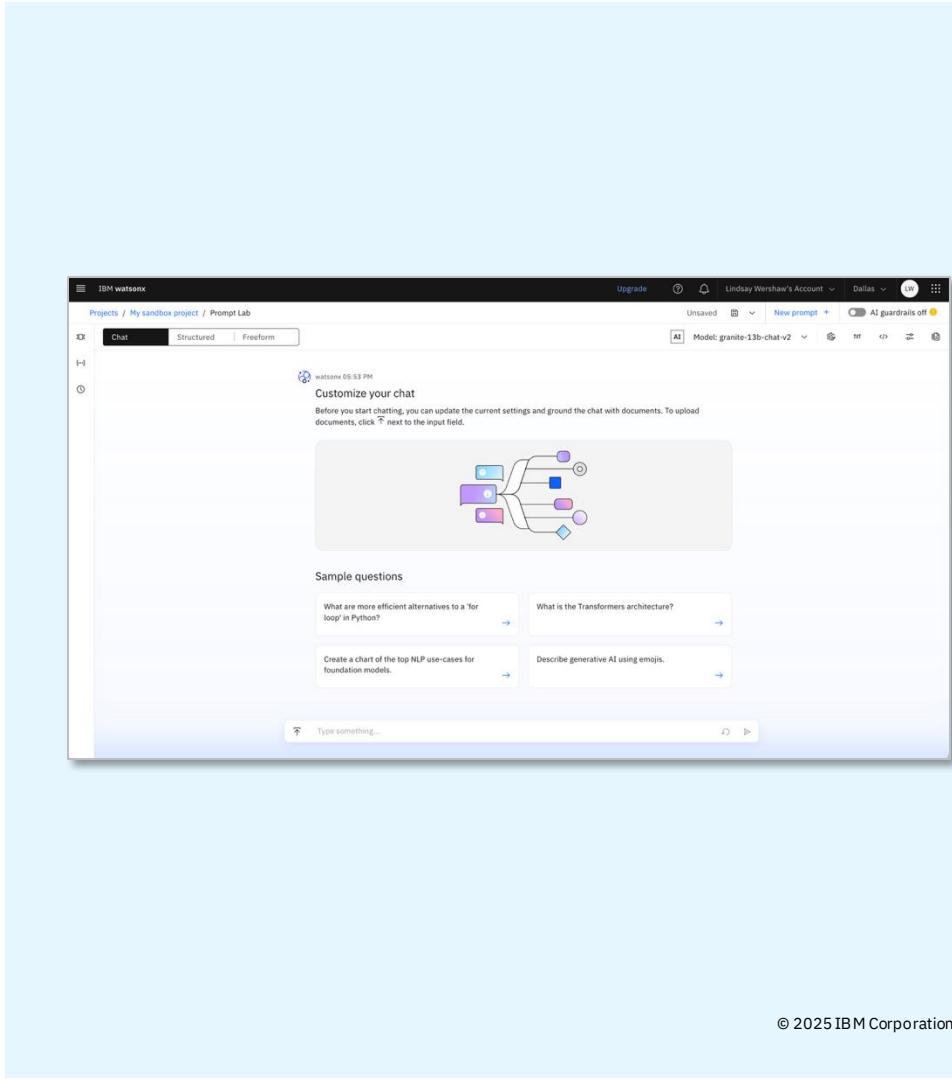
Prompt Lab

watsonx.ai Prompt Lab

Experiment with foundation models through an interactive user interface or an API.

- Easy to use chat, structured and free form prompt building interfaces
- Experiment with zero-shot, one-shot, or few-shot prompting to get the best results
- Adjustable AI guardrails
- Includes prompt examples for various use cases and tasks
- Save and share prompts for team collaboration
- Wide selection of foundation models to meet any task requirement
- Adjust model parameters to optimize results such as sampling, min/max tokens, stop sequences, repetition penalties, and more
- Export prompts and settings to a notebook to jump start development

Learn more in [documentation](#)



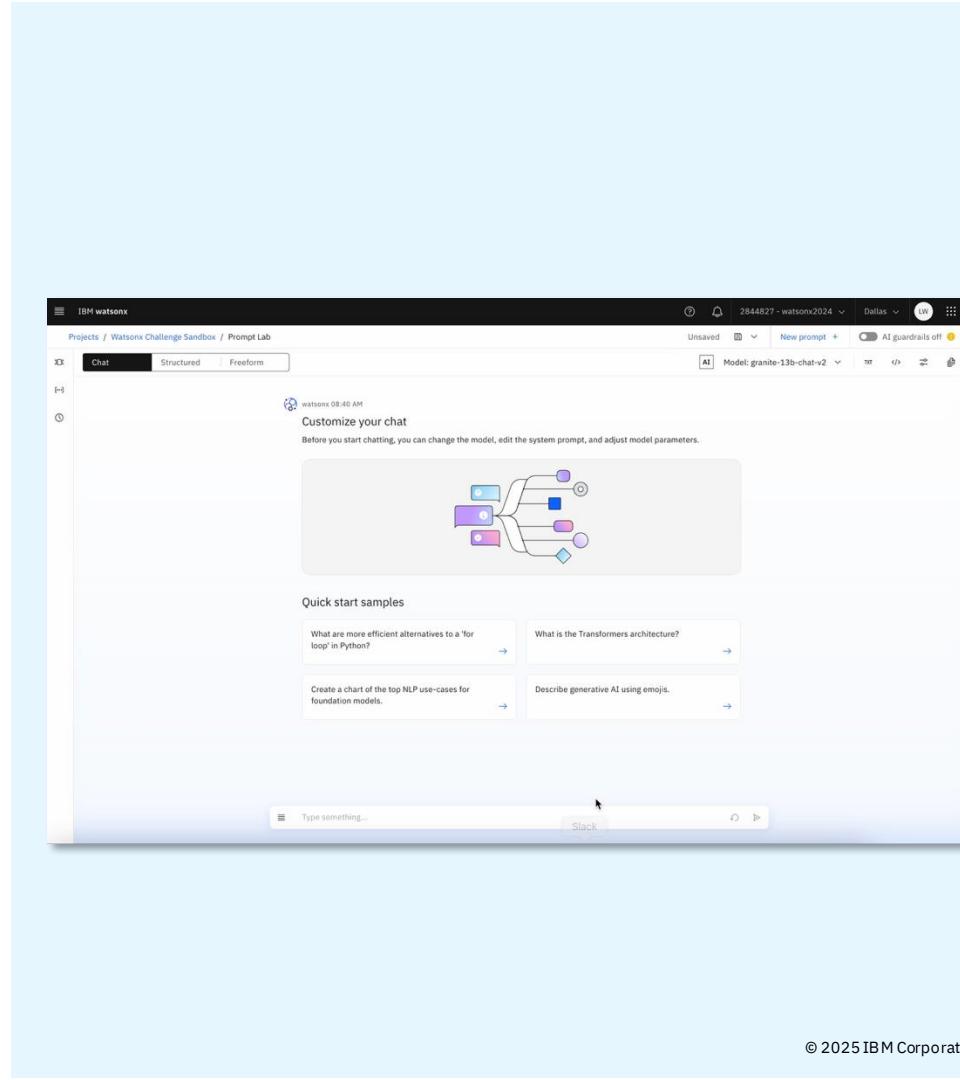
Chat with Documents

The Chat with Documents feature, within the Prompt Lab, provides an easy way to support RAG use cases.

Key benefits:

- Dramatically improve model output by grounding with relevant content
- Fast prototyping of your RAG use cases (e.g. test different LLMs, iterate on document index/collections, change embedding models, set chunk size and overlap, and more)
- Accelerate development by exporting the prompt as a python function deployed as a REST API
- Save prompts and settings as a template for team collaboration.

Learn more in [documentation](#)



Chat with Images

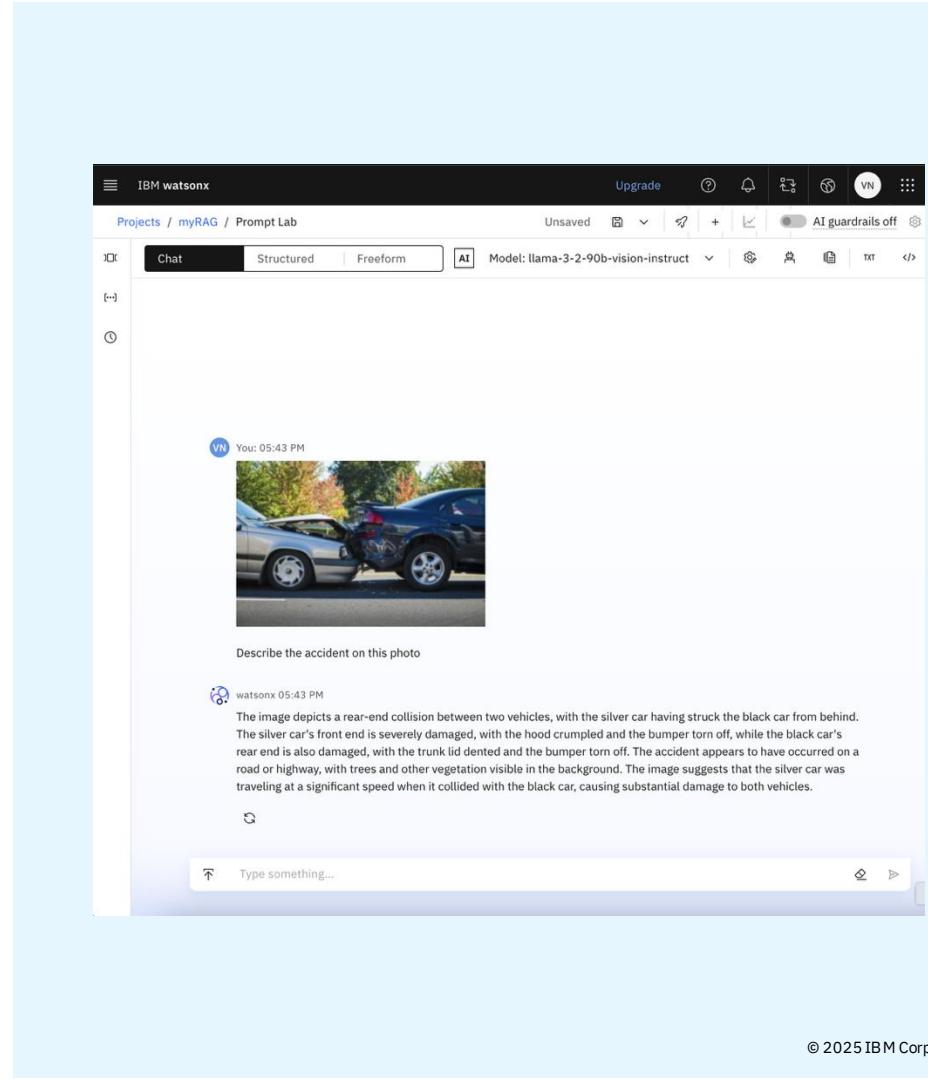
The Chat with Images feature, within the Prompt Lab, provides an easy way to convert visual information into text.

Example use cases:

- Automate the generation of alternative text for images to meet accessibility requirements
- Summarize photos to support use cases such as insurance claims (and many others)
- Convert images from a document into text before the document is used as grounding information for a RAG use case.

Chat with Images is supported by **Llama 3.2** (11b and 90b vision models).

Learn more in the [documentation](#) and [release blog](#).



RAG Solutions

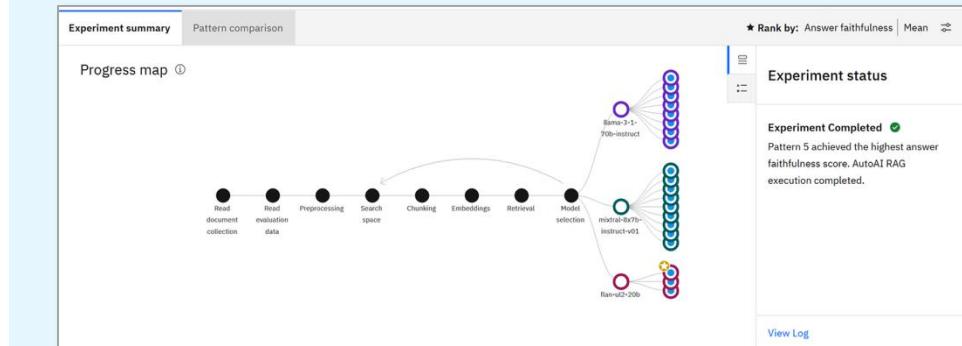
watsonX.ai AutoAI RAG

RAG prototypes are easy to build but hard to productionalize, which can require a team of experts and months of effort. AutoAI RAG helps to resolve this.

Key benefits:

- Accelerates design and deployment of optimized RAG systems based on client data and use case
- Quickly run various experiments to evaluate a constrained set of configuration options (models, chunk size, and more)
- Re-evaluate and modify recommended configurations when something changes (e.g. a new model version is released, quality of model responses change).

Learn more in [documentation](#).



AutoAI RAG tests a range of parameters (models, chunk size, etc.) in a series of experiments to automatically find the most optimal combination. It has 3 layers:

- Efficient RAG **hyper-parameter optimization** algorithm with end-to-end automation
- Best-of-breed RAG **evaluation metrics** and **benchmarking tools**
- **Parameterized RAG pipelines** for creating embeddings and for retrieval-based inference

watsonx.ai Text Extraction API

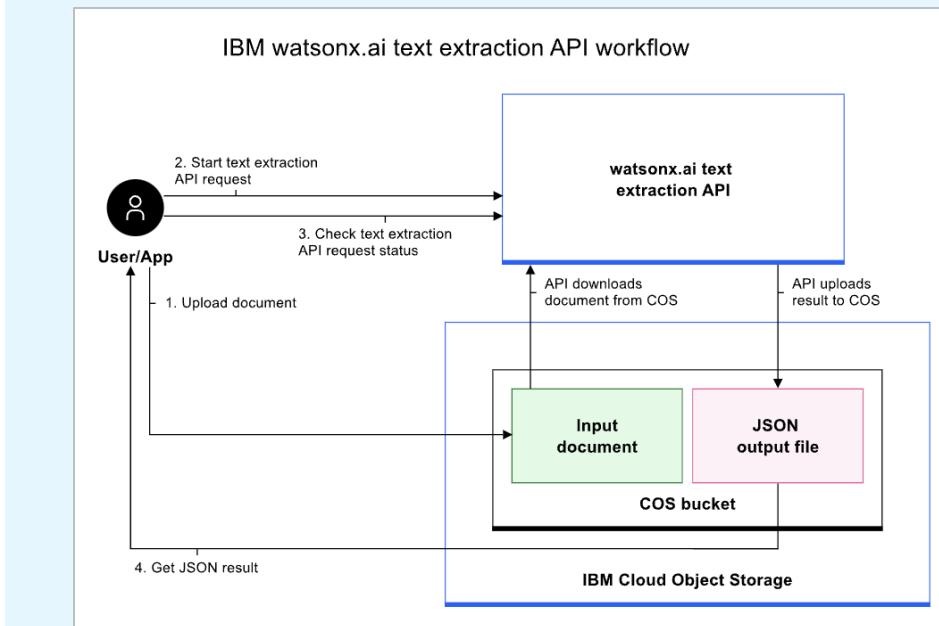
Extract content from documents which is essential for RAG use cases.

- Convert files with tables, diagrams and images into an AI-model friendly JSON file format
- Process the following file types: GIF, JPG, PDF, PNG, TIFF. Including scanned, hand-written documents
- Supports multiple languages

Technology:

- IBM's Natural Language Understanding (NLU) Service. See [Watson Document Understanding](#).
- Optical Character Recognition (OCR) to extract text from images.

Learn more in the [documentation](#).



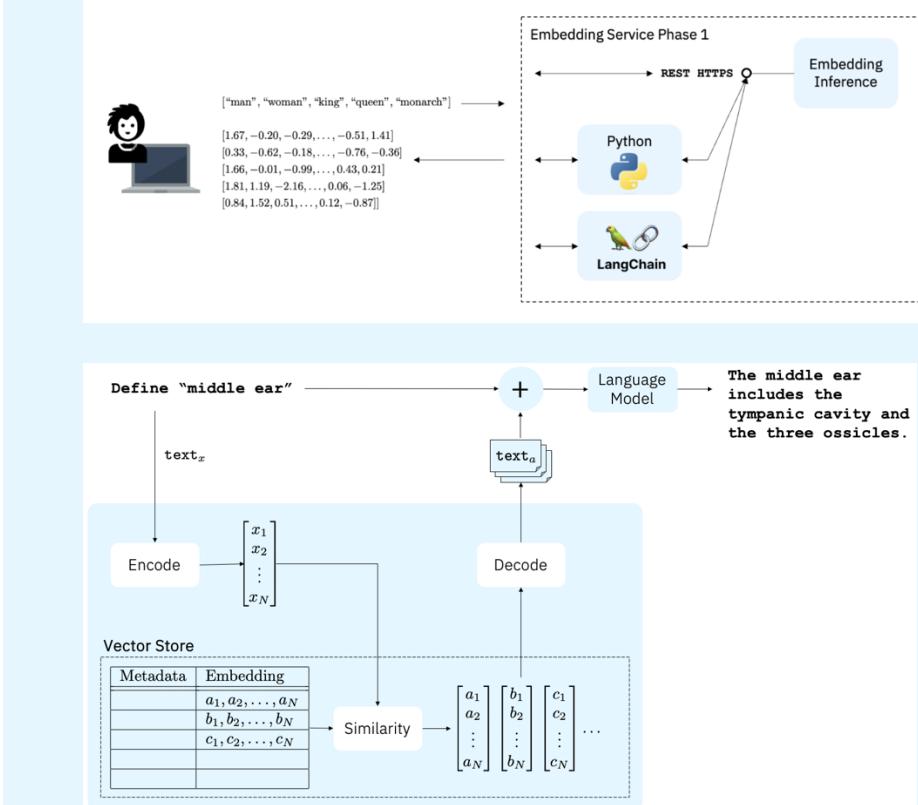
watsonx.ai Embeddings API

Convert text into dense vector embedding representations. Embeddings capture nuanced semantic and syntactic relationships between words and passages which are then stored in a vector database for retrieval.

Key benefits:

- Support RAG patterns with contextual grounding when utilizing embedding models for query and passage vectorization
- Improve RAG performance with semantically faithful representation of content, especially when compared to basic keyword-based alternatives in classic NLP modeling
- Efficient storage and compute profiles of embeddings make them easily infusible into generative AI applications

Learn more in [documentation](#). See list of [supported embedding models](#).



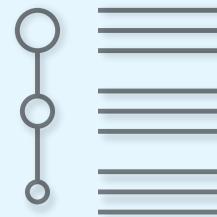
watsonx.ai Text Rerank API

Add precision to RAG use cases by improving the order of retrieved information with advanced re-ranking algorithms.

The Text Rerank API helps ensure relevant embeddings are used as context within a RAG pipeline, significantly enhancing performance.

- Reorder document passages (from most-to-least likely to answer) based on their similarity to a specified query
- Powered by the ms-marco-minilm-l-12-v2 reranker model ([see details](#))
- Fully integrated with the watsonx.ai REST API

Learn more in [documentation](#).

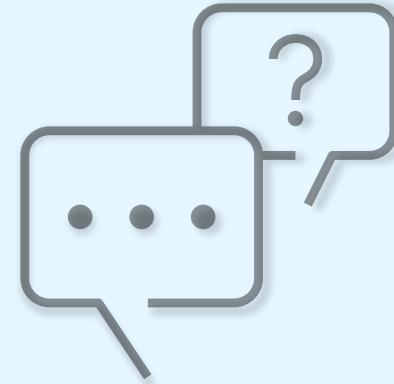


watsonx.ai Chat API

Create dynamic and adaptive conversational interfaces including chat-based application that require grounding documents, images, and tool calling for agent-driven applications.

- Build multi-user conversational workflows that use foundation models to generate answers.
- Easily identify different message types, such as a system prompt, user inputs, and foundation model outputs, including user-specific follow-up questions and answers.
- Supports granite, llama and mistral models (see [supported list](#))
- Fully integrated with the watsonx.ai API & SDK

Learn more about [adding chat functions](#) and [building agents](#) in documentation. See the [developer hub](#) for examples.



Agents

Three areas of agentic innovation

Pre-built agents

watsonx Orchestrate

Accelerate AI agent deployment.

Get started quickly with prebuilt AI agents powered with business logic and seamless integration to the tools that power your business.

Custom-built agents

watsonx.ai

Build custom-designed agents.

Design, deploy and manage AI agents with ease using pro- and low-code options.

Multi-agent orchestration

watsonx Orchestrate

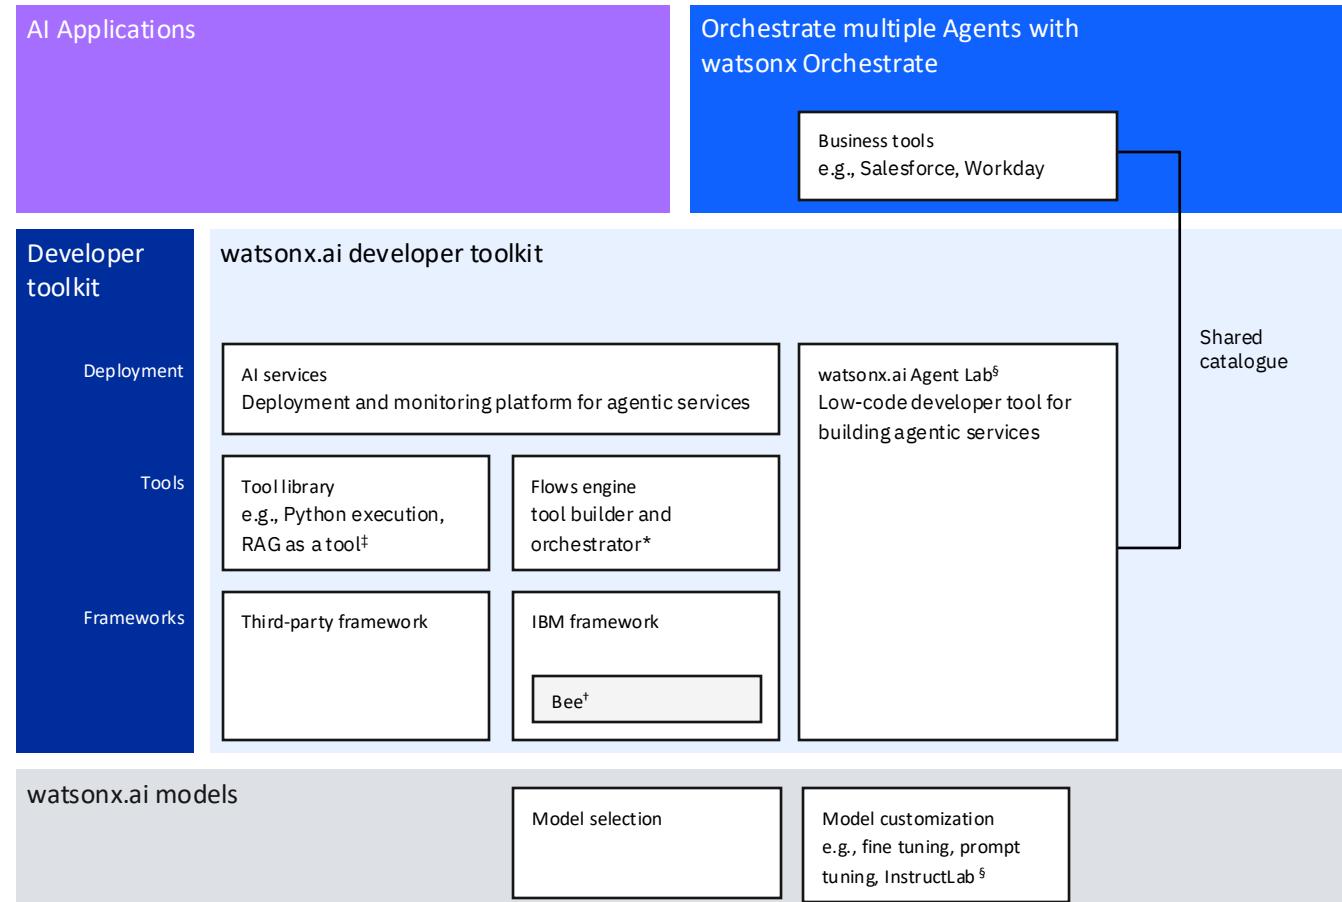
Manage all agents in one place.

Easily deploy and manage any agent for any task within a simple and unified user experience optimized to scale.

Architecture

Custom-built agents

One size does not fit all for building agents



* Stand-alone in tech preview; to be integrated into watsonx.ai
† Experimental, open-source project

‡ In discovery
§ Coming soon

watsonx.ai Agents

What's available today?

- Support for Chat API, Tool Calling, JSON returns together with providing the foundation for agentic support
- AI Services that provide the foundation for deployment of python based agentic solutions
- End-to-end template for the development and deployment of a custom Langraph agents
- Introduction of Agent Lab (Beta), enabling a low-code experience to build agentic solution with 1-click deployment as a production-ready API endpoint

The screenshot shows the IBM Watsonx interface for building an agent. On the left, the 'Build' tab is selected under the 'Projects / Kate's sandbox' header. The 'Model' dropdown is set to 'flan-ul2-20b'. The 'Agent preview' section on the right displays a 'Github Agent' entry from 12:34 pm, with a welcome message and instructions for creating GitHub issues. Below this, there are 'Quick start samples' for creating feature requests and bug templates. The main workspace shows configuration sections for 'Set up', 'Configuration', 'Framework' (LangGraph and React), and an 'Instruction' field with a placeholder for GitHub issues. A 'Tools' section lists three added tools: 'Web search' (retrieve info from the internet), 'Document search: Vector index name' (enable quick retrieval of information from vector indexes), and 'Python code executor' (enable dynamic code execution). A bottom search bar says 'Type something...'. The top navigation bar includes 'Upgrade', 'IBM account', 'Dallas', and other standard UI elements.

watsonx.ai SDKs

Build agentic services using popular open-source agentic frameworks through industry-standard API and SDK support.



LangGraph



LangChain



LlamaIndex



Bee

Lightweight Engine

watsonx.ai Lightweight Engine

Reduced footprint, inference-only installation
of watsonx.ai for client-managed software

Foundation Models

Access a variety of fit-for-purpose foundation models: IBM, open source, third party and imported custom foundation models.

Deploy Anywhere

Run watsonx.ai Lightweight Engine in any environment.

Reduced footprint

Run watsonx.ai Lightweight Engine with a smaller footprint.

Developer Experience

Leverage enterprise-grade APIs for a consistent developer experience across watsonx.ai full service and the lightweight engine.

watsonx.ai



Full service, enterprise-grade AI toolkit for building and deploying end-to-end AI solutions.

watsonx.ai
lightweight engine



lightweight engine to embed into external gen AI applications

watsonx.ai Lightweight Engine Features

Included

Features

- ✓ Foundation models (IBM, third-party, open-source)
- ✓ Imported custom foundation models
- ✓ watsonx.ai REST API
- ✓ watsonx.ai Text Embeddings API
- ✓ watsonx guardrails (HAP/PII filters)

Enterprise Capabilities

- | | |
|----------------------|-------------------------|
| ✓ Automated install | ✓ User Authentication |
| ✓ Horizontal Scaling | ✓ Logging |
| ✓ Backup and restore | ✓ Certification Manager |

Excluded, watsonx.ai full-service only

Features

- X Tuning Studio
- X Prompt Lab
- X Instruct Lab
- X Prompt Tuning via watsonx.ai REST API
- X Fine-tuning via watsonx.ai REST API¹
- X Chat with Documents and Images

Enterprise Capabilities

- X Projects and Spaces
- X Access Control

1 – on watsonx.ai full-service roadmap

Analyzing data and working with models

watsonx.ai Data Science and MLOps

Wide selection of integrated tooling to quickly build, deploy and manage machine learning models and AI applications

Build models

Comprehensive tooling for Model training and development

- Low/Pro code data science tools in a single environment
- Auto AI automates several aspects of the model lifecycle
- open-source libraries and integrated notebook-based interfaces
- Support for popular languages such as Python and R
- Integrated tooling for data pipelines and data preparation

Deploy models

Support for model deployment and management

- One-click deployment
- Continually deliver models into production
- Automated version control, retraining, code snippet generation for developers
- Build on OpenShift for reliability

Monitor models with watsonx.governance

Automate governance to drive trust across the AI lifecycle

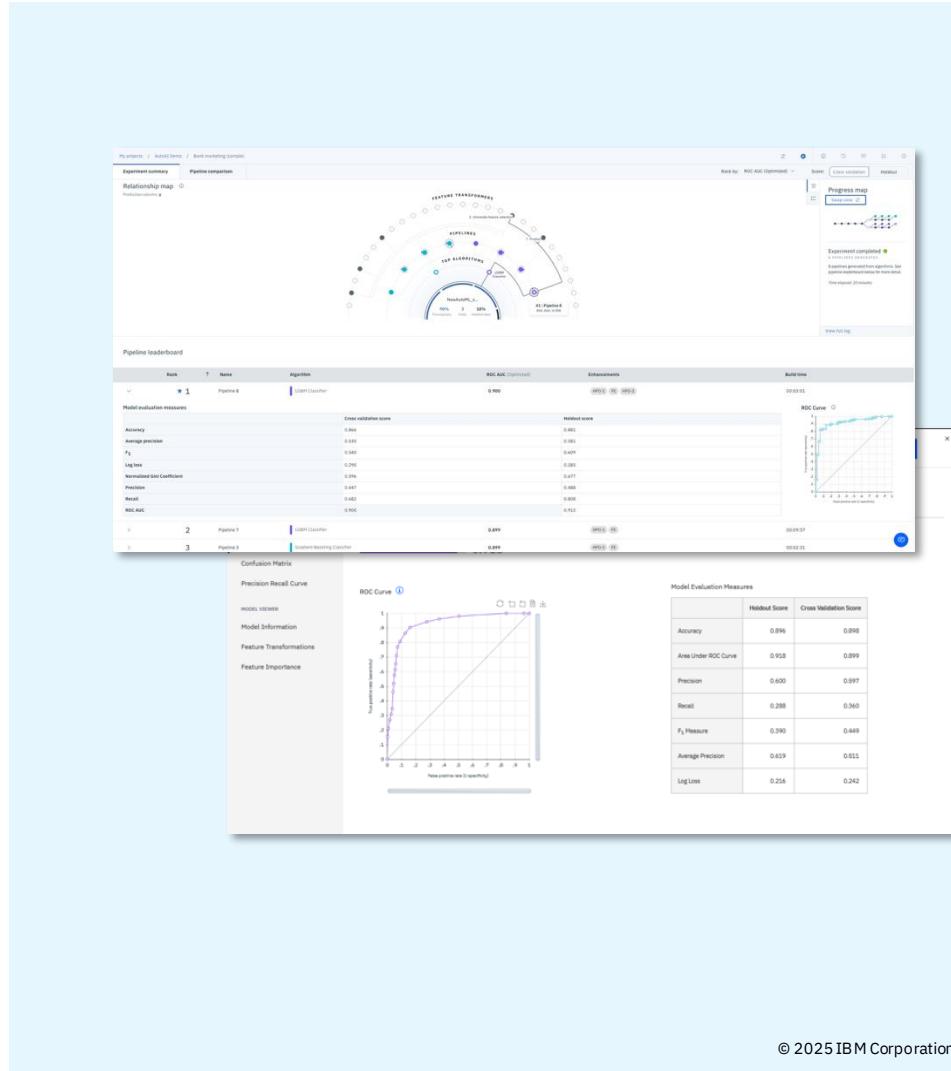
- Monitor drift over time, automatic retraining as required
- Provide model verification to stakeholders to ensure fairness, transparency and explainability over time
- Inform stakeholders of accuracy shifts before an issue occurs

watsonx.ai AutoAI

Expedite the creation of machine learning models with AutoAI, an easy to use, no-code interface. With a few clicks, quickly build powerful models such as classification, regression, time series forecasting, and more.

- Enhance the efficiency of the limited resources in a Data Science team by automating modeling.
- Quickly identify the best features, transforms, and models, as well as train the models, tune the hyperparameters, and rank the best-performing pipelines.
- Save an experiment as a notebook to understand the transformations applied to build the model or make changes.
- One click deployment of models, online or in batch mode

Learn more in [documentation](#)

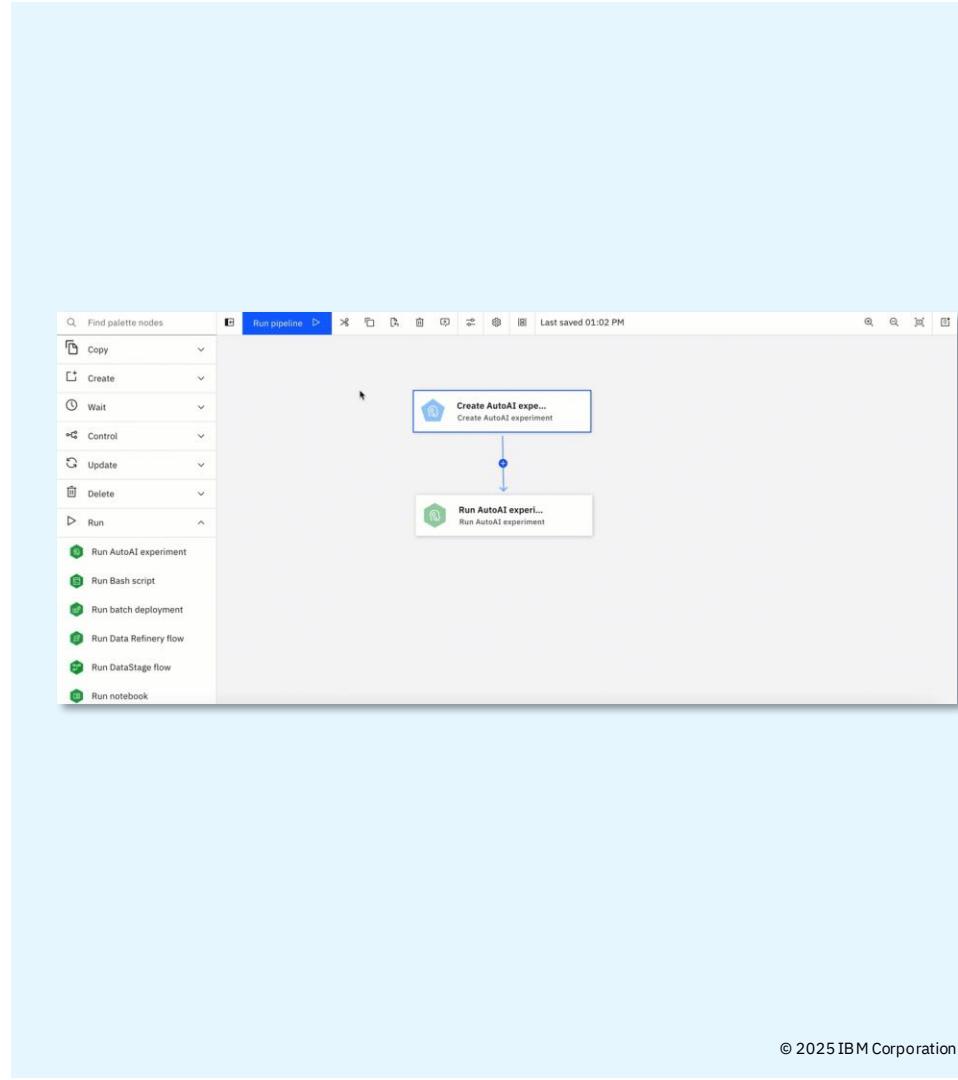


watsonx.ai Pipelines

Create end-to-end flow of assets from creation through deployment on a graphical canvas.

Key benefits:

- Increase efficiency and time savings with automated pipelines
- Combine tasks from different tools
 - Notebooks
 - Data refinery
 - Model creation
 - Model deployments
- Automate model training with new data sets
- Automatically deploy better challenger models



Learn more in [documentation](#)

watsonx.ai Synthetic Data Generator

Generate synthetic tabular data to address your data gaps

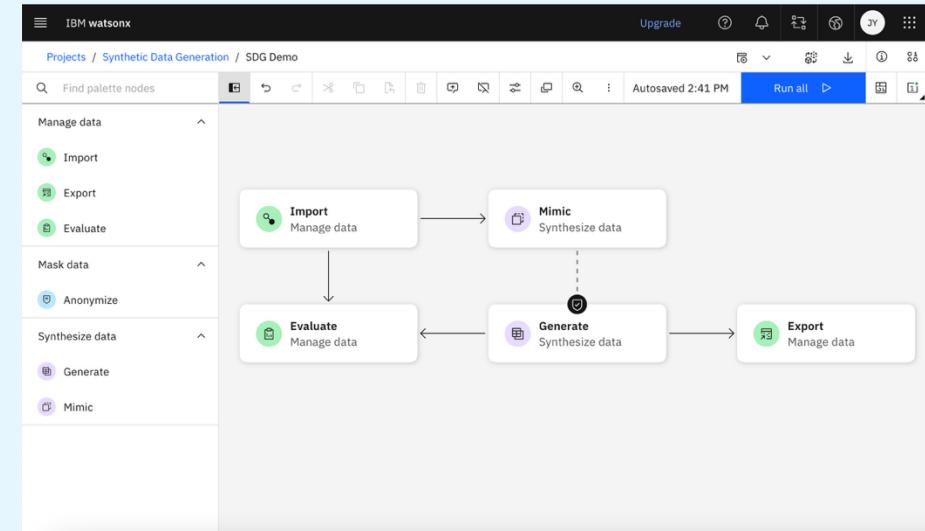
Create synthetic data at scale

- Unlock your valuable insights by using synthetic data.
- Create synthetic data using your existing data in a database or by uploading a file. Optionally, design your own data schema.
- Address data gaps and create synthetic edge cases to expedite classical AI model training.

Select your model & privacy needs

- Depending on your cost, fidelity, application, or data needs, you can select from multiple IBM models* to create your synthetic tabular data.
- When using existing data, IBM models apply differential privacy to minimize your privacy risk and give you control over the level of privacy protection required for your organization.

Learn more in [documentation](#)



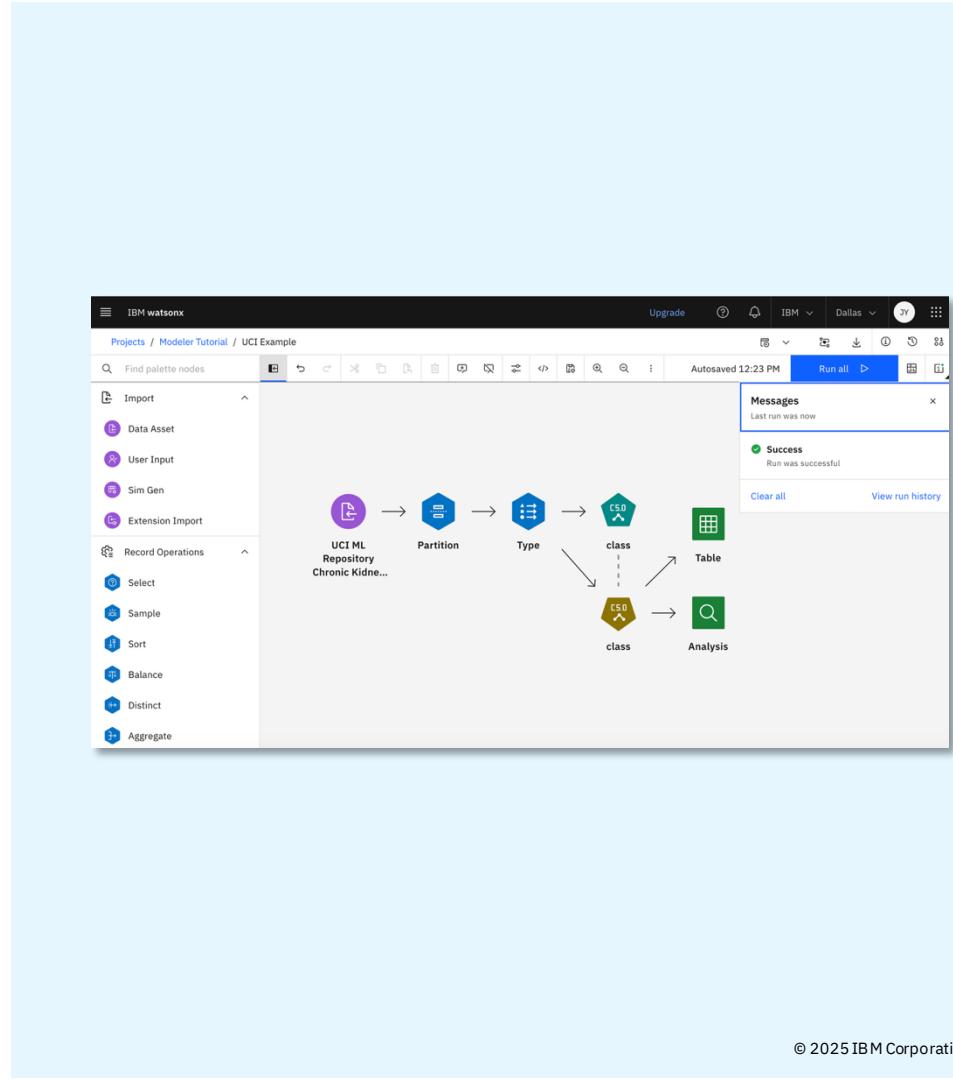
watsonx.ai Visual Modeling

The integrated SPSS Modeler brings a comprehensive set of nodes for data preparation, machine learning including auto modeling, data visualization, and push to production options.

Key benefits:

- Quickly build machine learning models using powerful drag and drop data analysis and visualization
- Drive direct insights from raw texts, uncover hidden value and patterns from text for prediction
- Deploy SPSS streams to into production

Learn more in [documentation](#)



watsonx.ai Decision Optimization

Leverage the power of prescriptive analytics with Decision Optimization (DO). Evaluate a problem holistically, identify constraints, and determine the best plan to accomplish your goal (e.g. max profit, min costs, optimal scheduling, etc.)

Key benefits:

- Solve optimization models by using DO's highly powerful mathematical and constraint programming engines
- Improve productivity by validating optimization models more quickly and easily using visual dashboards
- Create innovative solutions by combining optimization technology with data science techniques such as machine learning
- Operationalize your projects by deploying optimization models in production to drive business impact around real use cases

Learn more in [documentation](#)

The screenshot shows the IBM WatsonX interface with the title "IBM Watsonx" at the top. The main navigation bar includes "Upgrade", a help icon, a notification icon, a "Run configuration" button, a "Run" button, and other icons. The left sidebar is titled "Modeling Assistant" and lists several sections: "Build a scenario" (Scenario 1), "Prepare data" (1 table), "Build model" (Modeling Assistant, currently selected), "Explore solution", and "Analyze scenario data" (Visualization). The right side features four main cards: "Resource Assignment" (You want to assign (or match) resources to jobs, events, places.), "Scheduling" (You want to manage tasks, activities, events... with a precedence order.), "Selection & Allocation" (You want to select items (or combinations) from the same table.), and "Supply and Demand Planning" (You want to decide plant production, inventory and supply, per time period, to satisfy customer demand.). At the bottom, there is a note: "Not sure where to start? Use this question & answer guide." and "Don't see your type of decision? Code your model".

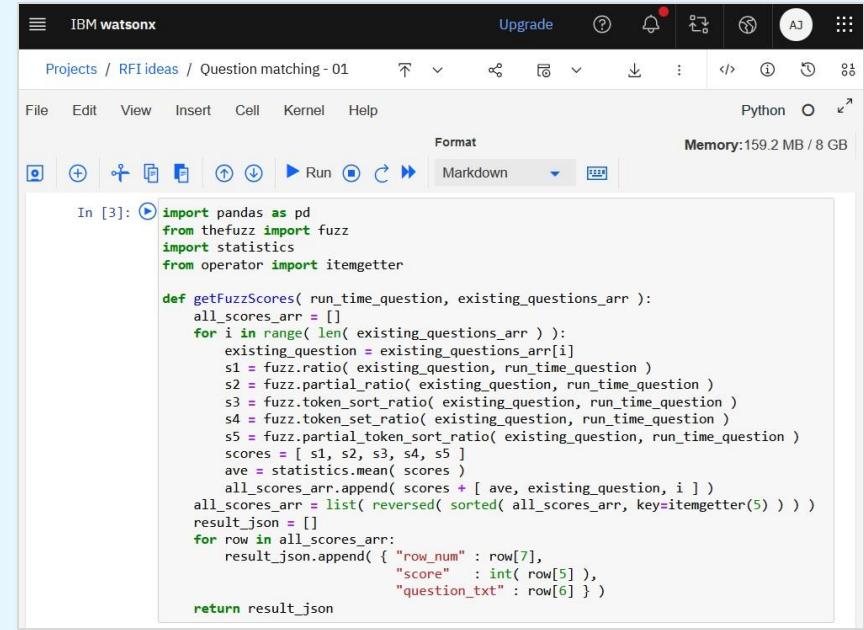
watsonx.ai Pro Code Tooling

Create, edit and execute Python and R code using Jupyter notebooks and scripts in code editors, for example the notebook editor or an the Rstudio integrated development environment (IDE)

Key benefits:

- Preinstalled IBM and opensource libraries
- Import your own libraries
- Connect notebooks to pipelines, publish to a catalog or Github for sharing and collaboration
- Automate repeatable processes through jobs to run the notebooks

Learn more in [documentation](#)



The screenshot shows the IBM Watsonx Pro Code Tooling interface. The top navigation bar includes 'IBM watsonx' (with a dropdown arrow), 'Upgrade', a user icon, and a gear icon. Below the bar are standard menu options: File, Edit, View, Insert, Cell, Kernel, and Help. To the right, it shows 'Python' and 'Memory: 159.2 MB / 8 GB'. The main area is a code editor with a toolbar above it containing icons for file operations, run, and refresh. A code cell labeled 'In [3]' contains the following Python script:

```
import pandas as pd
from thefuzz import fuzz
import statistics
from operator import itemgetter

def getFuzzScores( run_time_question, existing_questions_arr ):
    all_scores_arr = []
    for i in range( len( existing_questions_arr ) ):
        existing_question = existing_questions_arr[i]
        s1 = fuzz.ratio( existing_question, run_time_question )
        s2 = fuzz.partial_ratio( existing_question, run_time_question )
        s3 = fuzz.token_sort_ratio( existing_question, run_time_question )
        s4 = fuzz.token_set_ratio( existing_question, run_time_question )
        s5 = fuzz.partial_token_sort_ratio( existing_question, run_time_question )
        scores = [ s1, s2, s3, s4, s5 ]
        ave = statistics.mean( scores )
        all_scores_arr.append( scores + [ ave, existing_question, i ] )
    all_scores_arr = list( reversed( sorted( all_scores_arr, key=itemgetter(5) ) ) )
    result_json = []
    for row in all_scores_arr:
        result_json.append( { "row_num" : row[7],
                             "score" : int( row[5] ),
                             "question_txt" : row[6] } )
    return result_json
```

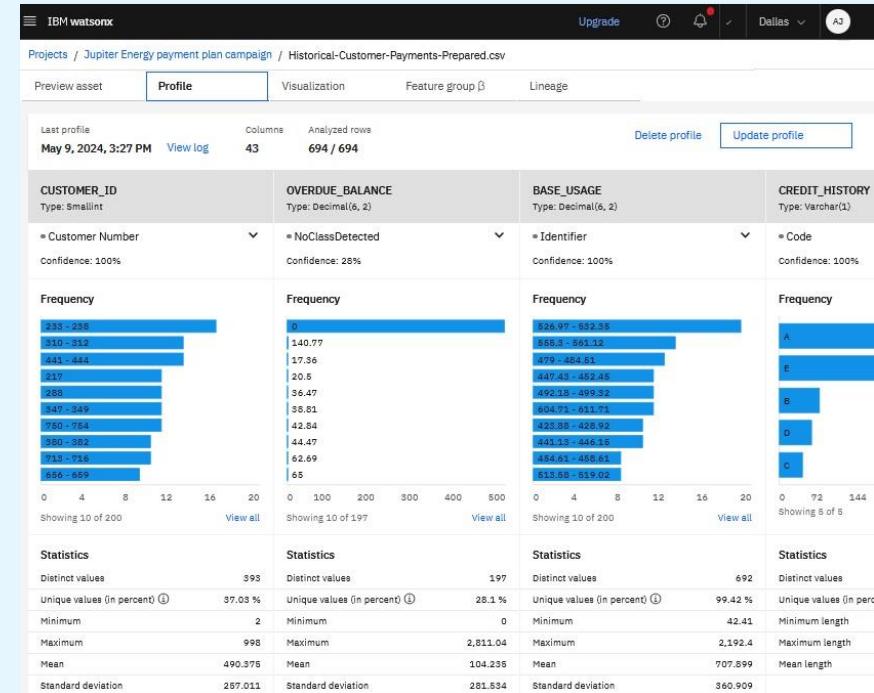
watsonx.ai Data Refinery

Conveniently analyze and prepare data for model training and all your AI application use cases within a single environment.

Key benefits:

- Retrieve or view tabular data from over 60 built-in data connectors
- Quickly analyze and visualize data with built-in data profiling and charting tools
- Cleanse and shape data with graphical editing with over 100 built-in operations
- Save time with scheduled flows and jobs for repeatable operations

Learn more in [documentation](#)



watsonx.ai Model Deployment

Expedite deployment of machine learning models, scripts, functions, and prompt templates for generative AI models.

- Secured with authentication credentials
- Support popular frameworks
- Deploy models in online or batch mode
- Score data through REST API calls
- Automate deployments with jobs

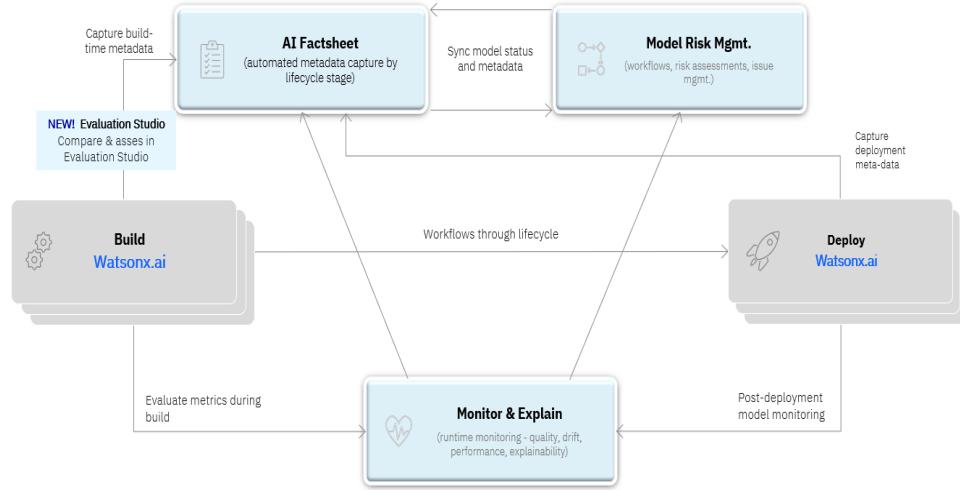


Learn more in [documentation](#)

watsonx.governance and watsonx.ai

Watsonx.governance brings **observability, explainability, and risk management** - key elements to build and deploy LLM applications and ML models.

- ✓ **Easily track and organize** AI Assets in Projects and Spaces by use case and associated lifecycle stage
- ✓ Enable AI Engineers to **easily determine the right Prompt asset** through experiments in Evaluation Studio
- ✓ **Evaluate quality and performance** of multiple task types (summarization, context generation, Q&A and more) during build with out-of-the-box metrics and integrated experience in Prompt Lab
- ✓ **Runtime monitoring** for quality, drift, and performance (and bias for ML models); user is alerted when thresholds are breached
- ✓ **RAG Root Cause Analysis insights** help AI engineers understand LLM responses
- ✓ **Assess hallucination** through quantifiable metrics for RAG task-type; **easily see** the phrases in **context used to generate responses**
- ✓ Reduce overhead of manual documentation through **automated documentation** of Prompt Template metadata and training input from the Prompt Lab

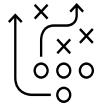


... and more!

- Out-of-the-box workflows through lifecycle
- Issue management integrated with use case
- Use case risk assessment
- Regulatory applicability assessments
- View AI use cases and deployments use across enterprise

Four ways to get started with watsonx.ai today

Looking forward to today's discussions



Free trial

Get started quickly by accessing AI models, data, and tools to build AI applications.



watsonx Developer Hub

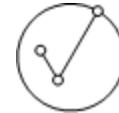
Access a library of templates and guides to get started building gen AI applications.



Request a client briefing or demo

Discussion and custom demonstration of IBM's generative AI watsonx point of view and capabilities. Understand where generative AI can be leveraged now for impact in your business.

2-4 hours onsite or virtual



5-year business value assessment

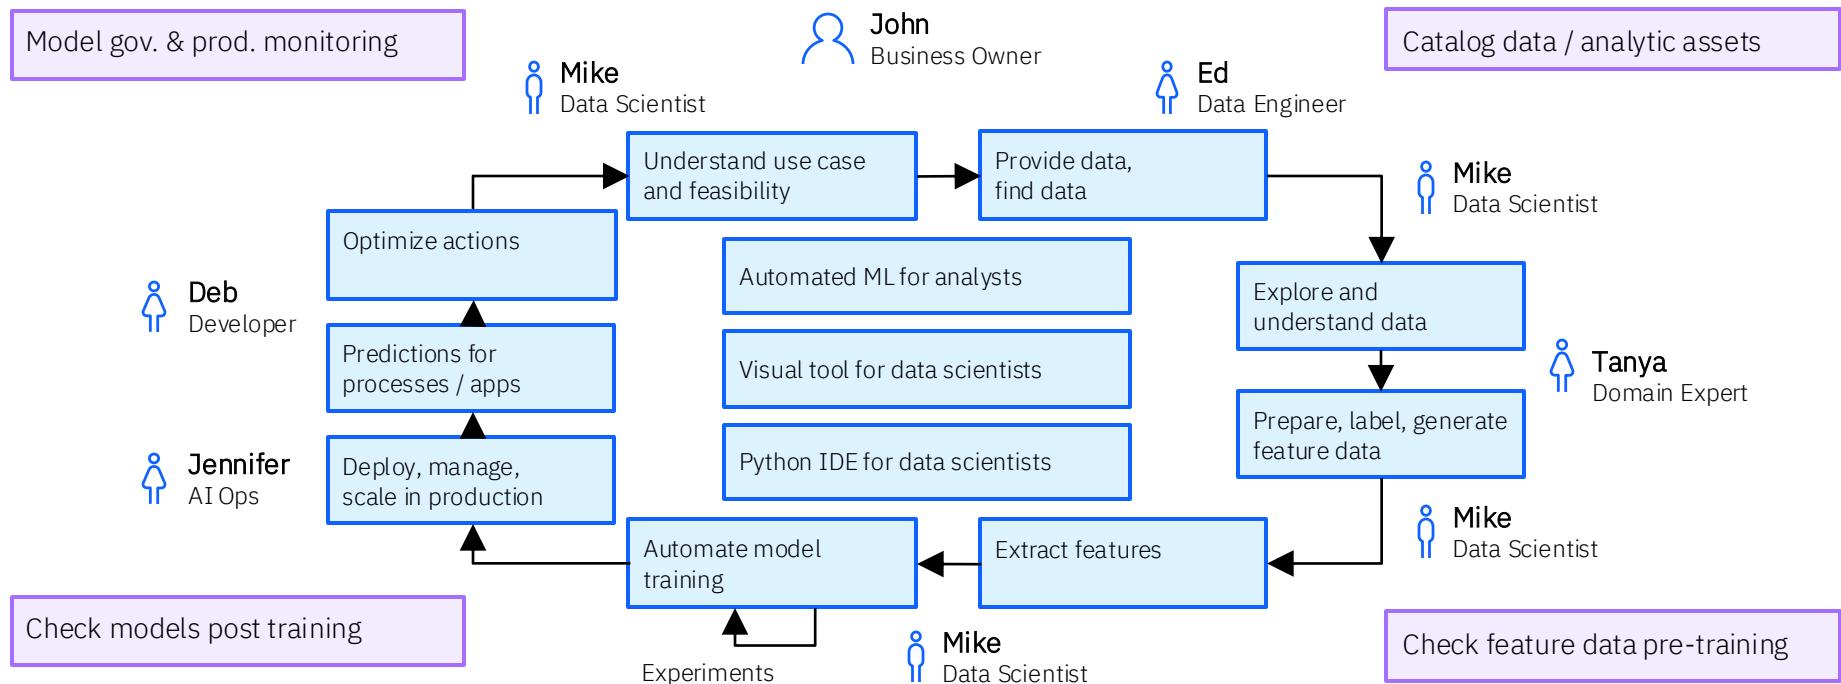
Engagement with an IBM multi-disciplinary team to jointly innovate and rapidly prove the business value of generative AI solutions using watsonx.

1-4 weeks



Backup

Data science requires a team



Model Customization Methods

Criteria	Prompt tuning	Parameter Efficient fine tuning (PEFT)	Full fine tuning	InstructLab
Technique	Prompts or embeddings are tuned while model parameters remain fixed.	Adapters, LoRA, QLoRA - Only a subset of parameters (e.g., adapters, biases) are fine-tuned.	SFTTrainer - All model parameters are fine-tuned on the target task.	Multi-phase aligned full fine tuning.
Test data reusability	Test data should not be reused; prompts are central	Test data should not be reused; prompts can guide which parameters are fine-tuned and can be reused.	Test data should not be reused; prompts can be incorporated into training data but may become less relevant as the model learns.	Reusable taxonomy across models and their versions
Model performance	Moderate performance; works well when the task is similar to the pre-trained knowledge.	High performance with reduced risk of overfitting; may not reach full fine-tuning levels.	Higher performance with sufficient data and compute resources.	Potentially the highest performance with less reliance on human data
Compute resources	Low; minimal compute required as only prompts or a small number of parameters are tuned.	Moderate; requires less compute than full fine-tuning due to selective parameter updates.	High computational and memory requirements due to the full update of model parameters.	High computational and memory requirements due to the full update of model parameters.
Time taken for tuning	Short; typically, the quickest tuning process as only prompts are optimized.	Moderate; faster than full fine-tuning but still requires time for training the selected parameters.	Long; depends on model size and dataset but generally requires significant time.	Long; depends on model size and synthetic data but generally requires significant time.
Summary	Most suitable for quick adaptation tasks, especially when computational resources are limited, or the task is closely related to the pre-trained model.	Ideal for resource-constrained environments or when needing to fine-tune multiple tasks efficiently.	Best for scenarios where maximum accuracy and task-specific adaptation are critical, and resources are abundant.	Much more than just a tuning technique. Doesn't need deep technical expertise and data prep. Provides a collaborative platform, synthetic data generation and better accuracy of models